

Dr. Emily Zhang

emily.zhang@email.com | (650) 555-0428 | Palo Alto, CA | linkedin.com/in/emilyzhang | github.com/ezhang

PROFESSIONAL SUMMARY

Distinguished AI Engineering Architect with 12 years of experience leading AI transformation initiatives at scale. Architected and deployed 25+ production AI agent systems serving 100M+ users globally. Recognized expert in agentic AI architecture, having spoken at major AI conferences and published influential research. Proven track record of building high-performing teams and delivering business impact through innovative AI solutions.

EXPERIENCE

Principal AI Architect / Head of AI Agents

Salesforce | January 2022 - Present

- Lead architecture and strategy for Salesforce AI Agent platform serving 150M+ users across Einstein GPT and Agentforce products
- Built team of 25+ AI engineers and researchers developing enterprise-grade autonomous agents
- Architected multi-agent orchestration system enabling complex business workflows with 95% reliability at scale
- Designed agent governance framework including safety guardrails, compliance controls, and audit trails
- Reduced customer service response time by 80% through intelligent routing and autonomous agent resolution
- Pioneered hybrid agent architecture combining symbolic reasoning with neural networks for explainable AI
- Established center of excellence for responsible AI development with ethics board and review processes
- Led integration of agents across Salesforce products (Sales Cloud, Service Cloud, Marketing Cloud)
- Filed 6 patents on agent architecture, tool orchestration, and adaptive learning systems

Director of AI Engineering

OpenAI | June 2018 - December 2021

- Led engineering team building GPT-3 API and early ChatGPT prototypes
- Architected scalable inference infrastructure serving 1M+ API requests per day
- Developed safety systems for content moderation and prompt injection prevention

- Built developer tools and SDKs enabling thousands of companies to integrate LLMs
- Collaborated with research team on RLHF implementation and alignment techniques
- Managed partnerships with enterprise customers on custom fine-tuning and deployment

Senior Engineering Manager

Google | July 2014 - May 2018

- Led Google Assistant AI team developing conversational agents for smart home and mobile
- Scaled conversation systems to support 500M+ devices across 30+ languages
- Built multi-modal agents integrating speech, vision, and text understanding
- Designed distributed training infrastructure for large-scale language models
- Mentored 15+ engineers and research scientists on ML system design

Research Scientist

DeepMind | August 2012 - June 2014

- Conducted fundamental research on multi-agent reinforcement learning
- Published 5 papers in top-tier venues (NeurIPS, ICML, ICLR)
- Contributed to AlphaGo project on search algorithms and policy networks

EDUCATION

Ph.D. in Artificial Intelligence

Stanford University

Dissertation: "Scalable Multi-Agent Reinforcement Learning Systems" | Published 8 papers | Graduated: 2014

Master of Science in Computer Science

Stanford University

Graduated: 2012

Bachelor of Science in Computer Science

Tsinghua University

Summa Cum Laude | Graduated: 2010

TECHNICAL SKILLS

Programming Languages: Python, C++, Go, TypeScript, Rust, Java, Scala

AI/ML Frameworks: PyTorch, TensorFlow, JAX, LangChain, LangGraph, LlamaIndex, AutoGen, Semantic Kernel, Ray, Horovod

LLM Platforms: OpenAI, Anthropic, Google Gemini, Cohere, Mistral, Together AI, Replicate

Agent Frameworks: n8n, AutoGPT, BabyAGI, CrewAI, AgentGPT, Superagent

Infrastructure: AWS, GCP, Azure, Kubernetes, Terraform, Istio, Ray Serve, Triton Inference Server

Data Systems: PostgreSQL, MongoDB, Redis, Kafka, Spark, Snowflake, Databricks

Vector Databases: Pinecone, Weaviate, Qdrant, Milvus, pgvector, Chroma

ML Ops: MLflow, Weights & Biases, LangSmith, Kubeflow, SageMaker, Vertex AI

Architecture: Microservices, Event-Driven Architecture, Distributed Systems, Real-time Processing

PROJECTS

Enterprise Agent Orchestration Platform

Designed and built production-grade platform for deploying, monitoring, and managing AI agents at scale. Handles 10M+ agent executions daily with sub-second latency and 99.99% uptime.

CERTIFICATIONS

- AWS Certified Solutions Architect - Professional
- Google Cloud Professional Cloud Architect
- Certified Kubernetes Application Developer (CKAD)